```
Traceback (most recent call last)
in <module>:4

  1 model.train()
  2 model = torch.compile(model, backend="hpu_backend")
  3 for epoch in range(1, 2):
❱ 4 |   train(epoch)
  5 |   acc = test()
  6 |   print(f'Epoch: {epoch}, Accuracy: {acc:.4f}')
  7

in train:8

   5 |   for i, (pos_rw, neg_rw) in enumerate(loader):
   6 |   |   optimizer.zero_grad()
   7 |   |   loss = model.loss(pos_rw.to(device), neg_rw.to(device))
❱  8 |   |   loss.backward()
   9 |   |   optimizer.step()
  10 |   |
  11 |   |   total_loss += loss.item()

/usr/local/lib/python3.10/dist-packages/torch/_tensor.py:531 in backward

   528 |   |   |   |   create_graph=create_graph,
   529 |   |   |   |   inputs=inputs,
   530 |   |   |   )
❱  531 |   |   torch.autograd.backward(
   532 |   |   |   self, gradient, retain_graph, create_graph, inputs=inputs
   533 |   |   )
   534

/usr/local/lib/python3.10/dist-packages/torch/autograd/__init__.py:289 in backward

  286 |   # The reason we repeat the same comment below is that
  287 |   # some Python versions print out the first line of a multi-line function
  288 |   # calls in the traceback and some print out the last line
❱ 289 |   _engine_run_backward(
  290 |   |   tensors,
  291 |   |   grad_tensors_,
  292 |   |   retain_graph,

/usr/local/lib/python3.10/dist-packages/torch/autograd/graph.py:768 in _engine_run_backward

  765 |   if attach_logging_hooks:
  766 |   |   unregister_hooks = _register_logging_hooks_on_whole_graph(t_outputs)
  767 |   try:
❱ 768 |   |   return Variable._execution_engine.run_backward(  # Calls into the C++ engine to
  769 |   |   |   t_outputs, *args, **kwargs
  770 |   |   )  # Calls into the C++ engine to run the backward pass
  771 |   finally:
```

NotImplementedError: Could not run 'aten::_sparse_coo_tensor_with_dims_and_tensors' with arguments from the 'SparseHPU' backend. This could be because the operator doesn't exist for this backend, or was omitted during the selective/custom build process (if using custom build). If you are a Facebook employee using PyTorch on mobile, please visit https://fburl.com/ptmfixes for possible resolutions. 'aten::_sparse_coo_tensor_with_dims_and_tensors' is only available for these backends: [HPU, Meta, SparseCPU, SparseMeta, BackendSelect, Python, FuncTorchDynamicLayerBackMode, Functionalize, Named, Conjugate, Negative, ZeroTensor, ADInplaceOrView, AutogradOther, AutogradCPU, AutogradCUDA, AutogradHIP, AutogradXLA, AutogradMPS, AutogradIPU, AutogradXPU, AutogradHPU, AutogradVE, AutogradLazy, AutogradMTIA, AutogradPrivateUse1, AutogradPrivateUse2, AutogradPrivateUse3, AutogradMeta, AutogradNestedTensor, Tracer, AutocastCPU, AutocastXPU, AutocastHPU, AutocastCUDA, FuncTorchBatched, BatchedNestedTensor, FuncTorchVmapMode, Batched, VmapMode, FuncTorchGradWrapper, PythonTLSSnapshot, FuncTorchDynamicLayerFrontMode, PreDispatch, PythonDispatcher].

HPU: registered at /npu-stack/pytorch-integration/hpu_ops/cpu_fallback.cpp:116 [backend fallback]
Meta: registered at /npu-stack/pytorch-fork/build/aten/src/ATen/RegisterMeta.cpp:26993 [kernel]
SparseCPU: registered at /npu-stack/pytorch-fork/build/aten/src/ATen/RegisterSparseCPU.cpp:1387 [kernel]
SparseMeta: registered at /npu-stack/pytorch-fork/build/aten/src/ATen/RegisterSparseMeta.cpp:287 [kernel]
BackendSelect: registered at /npu-stack/pytorch-fork/build/aten/src/ATen/RegisterBackendSelect.cpp:815 [kernel]
Python: registered at /npu-stack/pytorch-fork/aten/src/ATen/core/PythonFallbackKernel.cpp:153 [backend fallback]
FuncTorchDynamicLayerBackMode: registered at /npu-stack/pytorch-fork/aten/src/ATen/functorch/DynamicLayer.cpp:497 [backend fallback]
Functionalize: registered at /npu-stack/pytorch-fork/aten/src/ATen/FunctionalizeFallbackKernel.cpp:349 [backend fallback]
Named: registered at /npu-stack/pytorch-fork/aten/src/ATen/core/NamedRegistrations.cpp:7 [backend fallback]
Conjugate: registered at /npu-stack/pytorch-fork/aten/src/ATen/ConjugateFallback.cpp:17 [backend fallback]
Negative: registered at /npu-stack/pytorch-fork/aten/src/ATen/native/NegateFallback.cpp:18 [backend fallback]
ZeroTensor: registered at /npu-stack/pytorch-fork/aten/src/ATen/ZeroTensorFallback.cpp:86 [backend fallback]
ADInplaceOrView: fallthrough registered at /npu-stack/pytorch-fork/aten/src/ATen/core/VariableFallbackKernel.cpp:86 [backend fallback]
AutogradOther: registered at /npu-stack/pytorch-fork/torch/csrc/autograd/generated/VariableType_2.cpp:19857 [autograd kernel]
AutogradCPU: registered at /npu-stack/pytorch-fork/torch/csrc/autograd/generated/VariableType_2.cpp:19857 [autograd kernel]
AutogradCUDA: registered at /npu-stack/pytorch-fork/torch/csrc/autograd/generated/VariableType_2.cpp:19857 [autograd kernel]
AutogradHIP: registered at /npu-stack/pytorch-fork/torch/csrc/autograd/generated/VariableType_2.cpp:19857 [autograd kernel]
AutogradXLA: registered at /npu-stack/pytorch-fork/torch/csrc/autograd/generated/VariableType_2.cpp:19857 [autograd kernel]
AutogradMPS: registered at /npu-stack/pytorch-fork/torch/csrc/autograd/generated/VariableType_2.cpp:19857 [autograd kernel]
```
```

```
AutogradIPU: registered at /npu-stack/pytorch-fork/torch/csrc/autograd/generated/VariableType_2.cpp:19857 [autograd
kernel]
AutogradXPU: registered at /npu-stack/pytorch-fork/torch/csrc/autograd/generated/VariableType_2.cpp:19857 [autograd
kernel]
AutogradHPU: registered at /npu-stack/pytorch-fork/torch/csrc/autograd/generated/VariableType_2.cpp:19857 [autograd
kernel]
AutogradVE: registered at /npu-stack/pytorch-fork/torch/csrc/autograd/generated/VariableType_2.cpp:19857 [autograd
kernel]
AutogradLazy: registered at /npu-stack/pytorch-fork/torch/csrc/autograd/generated/VariableType_2.cpp:19857
[autograd kernel]
AutogradMTIA: registered at /npu-stack/pytorch-fork/torch/csrc/autograd/generated/VariableType_2.cpp:19857
[autograd kernel]
AutogradPrivateUse1: registered at /npu-stack/pytorch-fork/torch/csrc/autograd/generated/VariableType_2.cpp:19857
[autograd kernel]
AutogradPrivateUse2: registered at /npu-stack/pytorch-fork/torch/csrc/autograd/generated/VariableType_2.cpp:19857
[autograd kernel]
AutogradPrivateUse3: registered at /npu-stack/pytorch-fork/torch/csrc/autograd/generated/VariableType_2.cpp:19857
[autograd kernel]
AutogradMeta: registered at /npu-stack/pytorch-fork/torch/csrc/autograd/generated/VariableType_2.cpp:19857
[autograd kernel]
AutogradNestedTensor: registered at /npu-stack/pytorch-fork/torch/csrc/autograd/generated/VariableType_2.cpp:19857
[autograd kernel]
Tracer: registered at /npu-stack/pytorch-fork/torch/csrc/autograd/generated/TraceType_2.cpp:17623 [kernel]
AutocastCPU: fallthrough registered at /npu-stack/pytorch-fork/aten/src/ATen/autocast_mode.cpp:209 [backend
fallback]
AutocastXPU: fallthrough registered at /npu-stack/pytorch-fork/aten/src/ATen/autocast_mode.cpp:351 [backend
fallback]
AutocastHPU: fallthrough registered at
/builds/pytorch_modules_multi_build/torch/py3.10/pt2.4.0/Release/generated/backend/hpu_autocast_ops0.cpp:49
[backend fallback]
AutocastCUDA: fallthrough registered at /npu-stack/pytorch-fork/aten/src/ATen/autocast_mode.cpp:165 [backend
fallback]
FuncTorchBatched: registered at /npu-stack/pytorch-fork/aten/src/ATen/functorch/LegacyBatchingRegistrations.cpp:731
[backend fallback]
BatchedNestedTensor: registered at
/npu-stack/pytorch-fork/aten/src/ATen/functorch/LegacyBatchingRegistrations.cpp:758 [backend fallback]
FuncTorchVmapMode: fallthrough registered at
/npu-stack/pytorch-fork/aten/src/ATen/functorch/VmapModeRegistrations.cpp:27 [backend fallback]
Batched: registered at /npu-stack/pytorch-fork/aten/src/ATen/LegacyBatchingRegistrations.cpp:1075 [backend
fallback]
VmapMode: fallthrough registered at /npu-stack/pytorch-fork/aten/src/ATen/VmapModeRegistrations.cpp:33 [backend
fallback]
FuncTorchGradWrapper: registered at /npu-stack/pytorch-fork/aten/src/ATen/functorch/TensorWrapper.cpp:207 [backend
fallback]
PythonTLSSnapshot: registered at /npu-stack/pytorch-fork/aten/src/ATen/core/PythonFallbackKernel.cpp:161 [backend
fallback]
FuncTorchDynamicLayerFrontMode: registered at /npu-stack/pytorch-fork/aten/src/ATen/functorch/DynamicLayer.cpp:493
[backend fallback]
PreDispatch: registered at /npu-stack/pytorch-fork/aten/src/ATen/core/PythonFallbackKernel.cpp:165 [backend
fallback]
PythonDispatcher: registered at /npu-stack/pytorch-fork/aten/src/ATen/core/PythonFallbackKernel.cpp:157 [backend
fallback]
```